

4 Games: Adversarial Search

Until now, we considered the problems that are Single Agent, Fully Observable, Deterministic and Goal-oriented. There are many other problems involving multiple agents, sometimes adversarial to one another. Let us discuss the problems with multiple agents which are adversarial to each other.

Let us consider the scenario of a two player game. In this game, there are 2 bags and there are 2 balls in each bag with a number on it. There are 2 players, one aiming to maximize the score, let us call him/her as MAX player, and the other one is aiming to minimize the score, again let us call him/her as MIN player. The MAX player has to choose a bag, while the MIN player has to choose one ball out of the 2 balls in the chosen bag. The score is determined by the number on the ball chosen.

Let the 2 bags be B^1 and B^2 and the 2 balls in bag B^1 be B_1^1 and B_2^1 , and the 2 balls in bag B^2 be B_1^2 and B_2^2 . The Figure 9 shows the possible outcomes of the game. The red terminal in the Figure 9 represents the balls, and the number on them.




Figure 9: Example to Max-Min 2 player game.

As per Figure 9, s_0 , the root is the start state of this game, and s_3, s_4, s_5, s_6 as indicated by the red nodes are the terminal states where the score of the game is calculated. In addition, we have the 2 parameters depth d and branching factor b .

Let utility function be $Utility(s)$ where s is a terminal state, i.e., for example, $Utility(s_3) = 2$. The purpose of MAX player is to maximise the utility, while the purpose of MIN player is to minimize the utility.

An intuitive strategy that both players can adopt is to assume that the opponent would play optimally. That is, MAX player would always choose state with the maximum utility, and the MIN player would always choose the minimum utility. Let our transition model function be $RESULT(s, a)$. $RESULT(s, a)$ considers a state s and an action a taken as input to return the next state.

We can derive an algorithm, $MINIMAX(s)$ as follows:

$$MINIMAX(s) = \begin{cases} Utility(s) & \text{if } s \text{ is terminal state} \\ \max_{a \in A} MINIMAX(RESULT(s, a)) & \text{if MAX player} \\ \min_{a \in A} MINIMAX(RESULT(s, a)) & \text{if MIN player} \end{cases}$$

We will now trace this $MINIMAX$ algorithm on the tree in Figure 10 and 11. This will give us the tree below, where nodes are circled when they are explored.

- Max player will first explore the s_3 , and $Utility(s_3) = 2$, so at this point, we know that $Utility(s_1) \leq 2$.




Figure 10: Example to Max-Min 2 player game; Nodes s_3 and s_4 explored.

- As, Max player does not know the $Utility(s_1)$, it will explore s_4 , as $Utility(s_4) = 7$, as $Utility(s_1)$ will be decided by Min player, Max player know that the $Utility(s_1) = 2$. Refer to Figure 10.
- As shown in Figure 11, the next Max player will check for s_5 . $Utility(s_5) = 1$, so at this point, we know that $Utility(s_2) \leq 1$. From the prospective of Max player, terminal s_6 does not exists, as whatever be the Utility of s_6 , Max player knows that the $Utility(s_2) \leq 1$, as $Utility(s_1) = 2$, he should chose s_1 over s_2 .




Figure 11: Example to Max-Min 2 player game; Nodes s_3 , s_4 , s_5 explored.

Hence, Max player should chose B_1 in order to win the game.

4.1 Minimax Algorithm⁷

As discussed in two agent problem, the game can easily be solved by a recursive computation of the minimax values each node in the decision tree, working from the leaves to the root as the recursion unwinds. The algorithm is as follows assuming the MAX player starts first.

Algorithm 7 MinimaxDecision(S)

```
1:  $bestAction \leftarrow null$ 
2:  $max \leftarrow -\infty$ 
3: for all action  $a$  in  $ACTIONS(S)$  do
4:    $util \leftarrow MinValue(RESULT(S,a))$ 
5:   if  $util > max$  then
6:      $bestAction \leftarrow a$ 
7:      $max \leftarrow util$ 
8: return  $bestAction$ 
```

Algorithm 8 MaxValue(S)

```
1: if  $TERMINAL(S)=True$  then return  $UTILITY(S)$ 
2:  $v \leftarrow -\infty$ 
3: for all action  $a$  in  $ACTIONS(S)$  do
4:    $v \leftarrow \max(v, MinValue(RESULT(S,a)))$ 
5: return  $v$ 
```

Algorithm 9 MinValue(S)

```
1: if  $TERMINAL(S)=True$  then return  $UTILITY(S)$ 
2:  $v \leftarrow \infty$ 
3: for all action  $a$  in  $ACTIONS(S)$  do
4:    $v \leftarrow \min(v, MaxValue(RESULT(S,a)))$ 
5: return  $v$ 
```

The *MinimaxDecision* function can be modified to produce a set of optimal rational strategy for the game. A strategy is defined as a mapping of states in the game to an action or a series of actions. In other words, Strategy: $States \mapsto Actions$

4.1.1 Evaluation Time/Space Complexity of Minimax Algorithm

The *Minimax* algorithm implements a complete depth-first exploration of the game tree. The functions *MAXVALUE* and *MINVALUE* traverses the entire decision tree, down to each of the leaves, to evaluate the backed-up value of a state represented by a branch node by backward induction. Assuming each game state represented by a node has a branching factor b and the game tree has a maximum depth of m , the number of terminal states to evaluate would be b^m , which means the *Minimax* algorithm has a time complexity of $O(b^m)$. This is inefficient in practice. The space complexity of the *Minimax* algorithm is $O(bm)$ for if it generates the entire game tree at once, and $O(m)$ if it generates the actions of the game tree one at a time.

⁷Special thanks to Khiew Zhi Kai's class scribe for minimax algorithm, alpha-beta pruning and hill climbing notes.

4.2 Alpha-Beta Pruning

Next we move on to an alternative algorithm, Alpha-Beta Pruning, for the two player game that would optimize the *Minimax* algorithm, by ‘pruning’ off branches that cannot improve the outcome of the game for a player, so that not all terminal states have to be evaluated.

The α and β values are to be interpreted as such:

- α : Best utility (highest value) from MAX Player’s perspective
- β : Best utility (lowest value) from MIN Player’s perspective

Algorithm 10 MinimaxDecision(S)

```
1: bestAction  $\leftarrow$  null
2: max  $\leftarrow -\infty$ 
3: for all action  $a$  in ACTIONS( $S$ ) do
4:   util  $\leftarrow$  MinValue(RESULT( $S, a$ ),  $-\infty, \infty$ )
5:   if util > max then
6:     bestAction  $\leftarrow a$ 
7:     max  $\leftarrow$  util
8: return bestAction
```

Algorithm 11 MinValue(S, α, β)

```
1: if TERMINAL( $S$ )=True then return UTILITY( $S$ )
2:  $v \leftarrow \infty$ 
3: for all action  $a$  in ACTIONS( $S$ ) do
4:    $v \leftarrow \min(v, \text{MaxValue}(\text{RESULT}(S, a), \alpha, \beta))$ 
5:   if  $v \leq \alpha$  then return  $v$ 
6:    $\beta \leftarrow \min(v, \beta)$ 
7: return  $v$ 
```

Algorithm 12 MaxValue(S, α, β)

```
1: if TERMINAL( $S$ )=True then return UTILITY( $S$ )
2:  $v \leftarrow -\infty$ 
3: for all action  $a$  in ACTIONS( $S$ ) do
4:    $v \leftarrow \max(v, \text{MinValue}(\text{RESULT}(S, a), \alpha, \beta))$ 
5:   if  $v \geq \beta$  then return  $v$  ▷ Pruning of branch
6:    $\alpha \leftarrow \max(\alpha, v)$ 
7: return  $v$ 
```

An alternative call to the algorithm to determine the utility of the MAX player if both players play rationally is MAXVAL($root, -\infty, \infty$).

4.2.1 An Example for Alpha-Beta Pruning

Let us consider the game tree in Figure 12, where MAX player begins first. We will now trace the algorithm’s logic as it processes the game tree, as well as how the algorithm avoids having to evaluate

some terminal states.

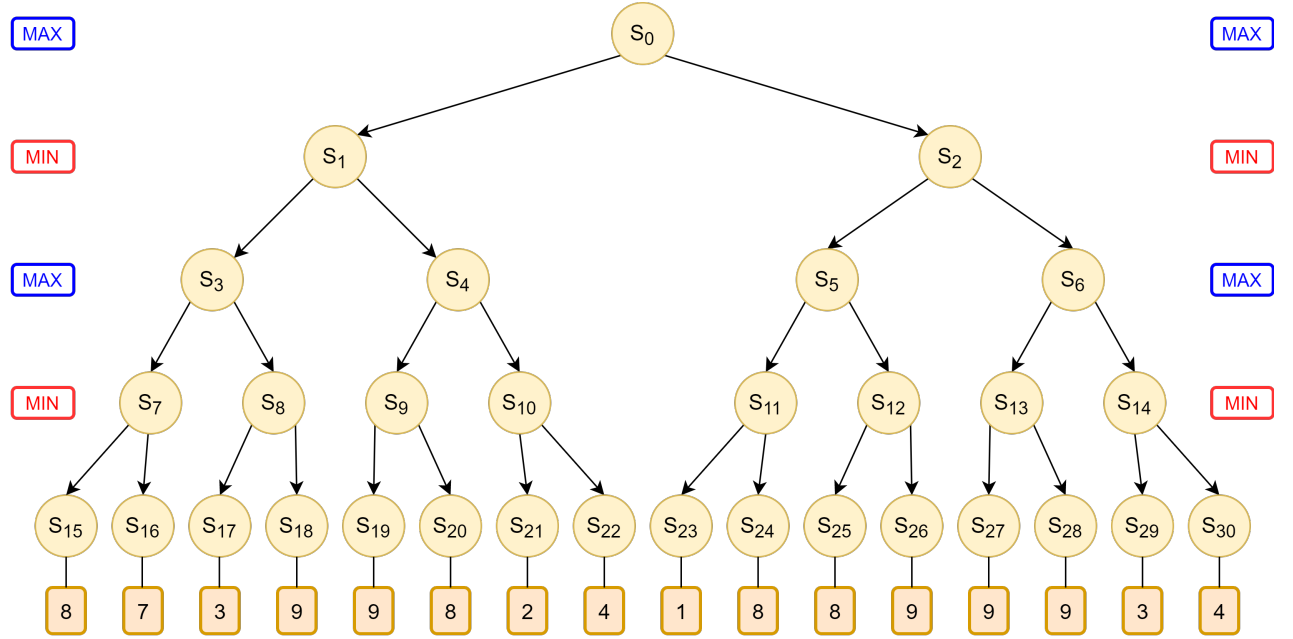


Figure 12: Game tree where MAX Player starts first

The game tree in Figure 13 is annotated to reflect the logic of the Alpha-Beta Pruning.

The algorithm evaluates the nodes from left to right along the game tree recursively. The logic flows in the following order, denoted by the increasing numbers on the white boxes in Figure 13:

- **Box 1:** The algorithm recursively traverses down from root node S_0 to S_{15} to evaluate the first terminal state, which results in a utility of 8
- **Box 1.1:** The MIN player at S_7 , knowing that the terminal state at S_{15} will produce utility of 8, now knows that the utility at S_7 will be equal to or less than 8, because he could potentially choose S_{16} which might have a lower utility.
- **Box 2:** The algorithm evaluates the next terminal state S_{16} to find a utility of 7.
- **Box 2.1:** For the MIN player at S_7 , the utility of 7 at S_{16} is lower than utility of 8 at S_{15} . Thus, he will definitely choose S_{16} , and utility at S_7 is assigned as 7.
- **Box 2.2:** The MAX player at S_3 , knowing that choosing S_7 will result in a utility of 7, now knows that his utility will be at least 7, because he can potentially choose S_8 that could have a higher utility value.
- **Box 3:** The algorithm evaluates the next terminal state S_{17} to find a utility of 3.
- **Box 3.1:** The MIN player at S_8 , now knows that his utility will be at most 3 from the terminal state S_{17} , because he can potentially choose S_{18} which could have a lower utility.
- **Box 3.2:** The MAX player at S_3 now knows that by choosing S_8 , his utility will be at most 3, which is worse than the option of choosing S_7 that gives him utility 7. Thus he will choose S_7 , and the entire branch to S_8 is pruned off, without having to evaluate the remaining nodes.

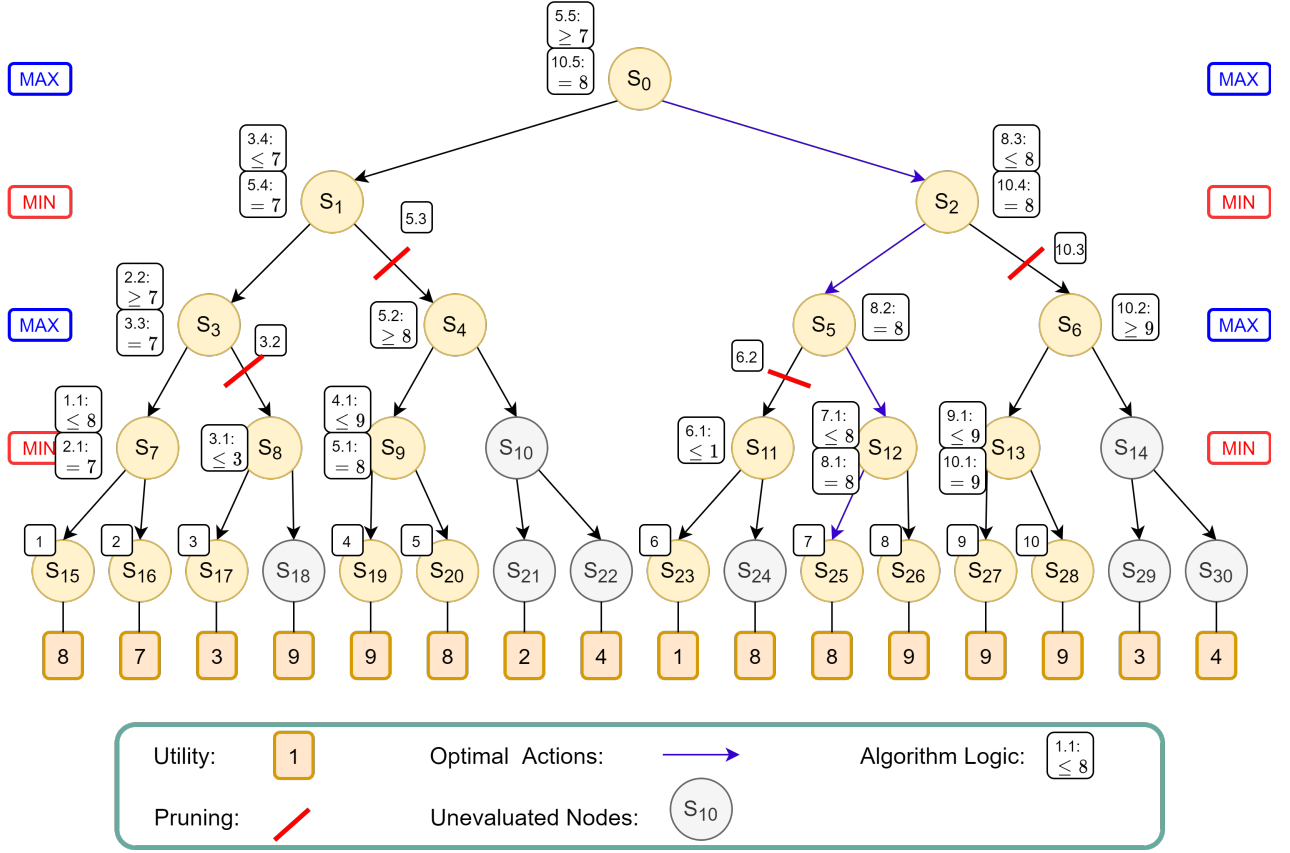


Figure 13: Game tree annotated with Alpha-Beta Pruning Logic

- **Box 3.3:** With only S_7 remaining, a utility value of 7 is assigned to node S_3 .
- **Box 3.4:** The MIN player at S_1 , knowing that choosing S_3 will result in a utility of 7, now knows that the utility at S_1 is at most 7, because he can potentially choose S_4 which could have a lower utility.
- **Box 4:** The algorithm evaluates the next terminal state S_{19} to find a utility of 9.
- **Box 4.1:** The MIN player at S_9 , now knows that his utility will be at most 9 from the terminal state S_{19} , because he can potentially choose S_{20} which could have a lower utility.
- **Box 5:** The algorithm evaluates the next terminal state S_{20} to find a utility of 8.
- **Box 5.1:** The MIN player at S_9 would choose S_{20} which gives a lower utility of 8 compared to S_{19} which gives a utility of 9. Thus the lower utility of 8 is assigned to S_9 .
- **Box 5.2:** The MAX player at S_4 , knowing that choosing S_9 results in a utility of 8, knows that his at S_4 utility is at least 8 because he could potentially choose S_{10} which might have a higher utility.

- **Box 5.3:** The MIN player at S_1 knows it can pick S_3 for utility 7, which is better than picking S_4 which grants the utility of at least 8. Therefore, S_4 will not be picked, at the entire branch of S_4 is pruned.
- **Box 5.4:** The MIN player at S_1 will therefore pick S_3 to guarantee a lower utility of 7, thus utility value 7 is assigned to S_1 .
- **Box 5.5:** The MAX player at S_0 , knowing that picking S_1 results in a utility of 7, knows that at S_0 , his utility is at least 7, because he could potentially pick S_2 which may have a higher utility.
- **Box 6:** The algorithm evaluates the next terminal state S_{23} to find a utility of 1.
- **Box 6.1:** The MIN player at S_{11} , knowing that picking S_{23} results in utility 1, now knows that his utility at S_{11} is at most 1, because he can potentially choose S_{24} which may yield a lower utility.
- **Box 6.2:** The MAX player at S_5 , knowing that choosing S_{11} will result in a utility at most 1, will not choose S_{11} , because he could have chosen S_1 when he was at S_0 and get utility 7. Thus the entire branch of S_{11} is pruned. The utility at S_5 is still uncertain as choosing S_{12} could still potentially yield a higher utility than 7.
- **Box 7:** The algorithm evaluates the next terminal state S_{25} to find a utility of 8.
- **Box 7.1:** The MIN player at S_{12} , knowing that choosing S_{25} results in a utility of 8, now knows that his utility at S_{12} is at most 8, because he can potentially choose S_{26} which might have a lower utility.
- **Box 8:** The algorithm evaluates the next terminal state S_{26} to find a utility of 9.
- **Box 8.1:** The MIN player at S_{12} will choose S_{25} which provides a lower utility of 8. Thus S_{12} is assigned the value of 8.
- **Box 8.2:** The MAX player at S_5 now knows that choosing S_{12} returns a utility of 8, which is higher than the utility of 7 had he chosen to choose S_1 when he was at S_0 . Thus, a value of 8 is assigned to S_5 .
- **Box 8.3:** The MIN player at S_2 knows that choosing S_5 results a utility of 8, thus he now knows that the utility at S_2 is at most 8, because he can potentially choose S_6 which may potentially yield a lower utility.
- **Box 9:** The algorithm evaluates the next terminal state S_{27} to find a utility of 9.
- **Box 9.1:** The MIN player at S_{13} knows that choosing S_{27} results in a utility of 9, and now knows that his utility at S_{13} is at most 9, because he can potentially choose S_{28} which may yield a lower utility.
- **Box 10:** The algorithm evaluates the next terminal state S_{28} to find a utility of 9.
- **Box 10.1:** The MIN player at S_{13} knows that choosing either S_{27} or S_{28} results in a utility of 9, thus assigning a utility of 9 for S_{13} .
- **Box 10.2:** The MAX player at S_6 knows that choosing S_{13} results in a utility of 9, thus his utility at S_6 is at least 9, because he can potentially choose S_{14} which may potentially yield a higher utility.

- **Box 10.3:** The MIN player at S_2 knows that choosing S_6 results in a utility of at least 9, which is worse than choosing S_5 which guarantees a lower utility of 8. Thus, the MIN player will not choose S_6 , and the entire branch of S_2 is pruned.
- **Box 10.4:** The MIN player at S_2 will therefore choose S_5 to get a utility of 8. Thus S_2 is assigned utility of 8.
- **Box 10.5:** The MAX player at S_0 knows that choosing S_2 will result in a utility of 8, which is higher than the utility of 7 by choosing S_1 . Thus, the MAX player will choose S_2 to get a utility of 8 at S_0 .

4.2.2 Evaluation of Time/Space Complexity of Alpha-Beta Pruning

One weakness of the Alpha-beta pruning algorithm is, it is sensitive to the order at which the terminal states are evaluated. If the algorithm happens to evaluate the worse successors first in all nodes, there would be no chance for optimization by pruning. Therefore in the worse case, the algorithm runs the same as the Minimax Search at $O(b^m)$ in time complexity. However, if the algorithm knows where to priorities the search of the best successor nodes first, it only need to examine $b^{m/2}$ nodes, and therefore improves its time complexity to $O(b^{m/2})$ in the best case. If the successors are evaluated randomly, the algorithm's time complexity becomes $O(b^{3m/4})$ on average. The space complexity is similar to that of the minimax algorithm.

4.3 Heuristic Function for Imperfect Decision-Making

Because Alpha-beta pruning still has to explore the game tree until the terminal state, this makes its time complexity impractical in games with large depth such as chess. One strategy is to cut off the search earlier, and apply a heuristic evaluation function to each state at the cutoff, thus allowing for non-terminal nodes to become leaves of the smaller search tree.

A good heuristic evaluation function should have the following properties:

1. The evaluation function should evaluate terminal states in a similar preference order as the true utility function would on the terminal states.
2. Computation of the evaluation function should not take too long (ideally close to constant time complexity)
3. For non-terminal states, the evaluation function should choose states highly correlated with the chances of winning. The uncertainty is induced by computational limits.

For example in chess, an evaluation function the expected probability of a win based on a weighted linear function of certain features of the game. For example, the evaluation function can be,

$$EVAL(S) = w_1 f_1(S) + w_2 f_2(S) + \dots + w_n f_n(S)$$

where w_i refers to the weights, while $f_i(S)$ refers to a computation of features in the game state, such as the number of a type of chess piece on the board. The evaluation function need not calculate the actual probability of winning, but the ordinal relation between states via the evaluation function should be similar to that of the expected probability of winning.

The Alpha-beta pruning algorithm could implement the cut-off and heuristic by modifying the terminal condition of the algorithm.

Algorithm 13 MinimaxDecision(S)

```
1:  $bestAction \leftarrow null$ 
2:  $max \leftarrow -\infty$ 
3: for all action  $a$  in  $ACTIONS(S)$  do
4:    $util \leftarrow MinValue(RESULT(S,a), 1, -\infty, \infty)$ 
5:   if  $util > max$  then
6:      $bestAction \leftarrow a$ 
7:      $max \leftarrow util$ 
8: return  $bestAction$ 
```

Algorithm 14 MinValue($S, depth, \alpha, \beta$)

```
1: if CUTOFFTEST( $S, depth$ )=True then return EVAL( $S$ )
2:  $v \leftarrow \infty$ 
3: for all action  $a$  in  $ACTIONS(S)$  do
4:    $v \leftarrow \min(v, MaxValue(RESULT(S,a), depth + 1, \alpha, \beta))$ 
5:   if  $v \leq \alpha$  then return  $v$ 
6:    $\beta \leftarrow \min(v, \beta)$ 
7: return  $v$ 
```

Algorithm 15 MaxValue($S, depth, \alpha, \beta$)

```
1: if CUTOFFTEST( $S, depth$ )=True then return EVAL( $S$ )
2:  $v \leftarrow -\infty$ 
3: for all action  $a$  in  $ACTIONS(S)$  do
4:    $v \leftarrow \max(v, MinValue(RESULT(S,a), depth + 1, \alpha, \beta))$ 
5:   if  $v \geq \beta$  then return  $v$  ▷ Pruning of branch
6:    $\alpha \leftarrow \max(\alpha, v)$ 
7: return  $v$ 
```

One way to control the amount of search is to set a depth limit d so that CUTOFFTEST($S, depth$) returns true when depth reaches limit d or a terminal state is reached whichever is earlier. The depth limit can be determined by the deepest depth a computer can explore while just keeping to the time limit required to make a decision. With a cut-off, the time complexity improves to $O(b^{3d/4})$ on average where $d < m$.